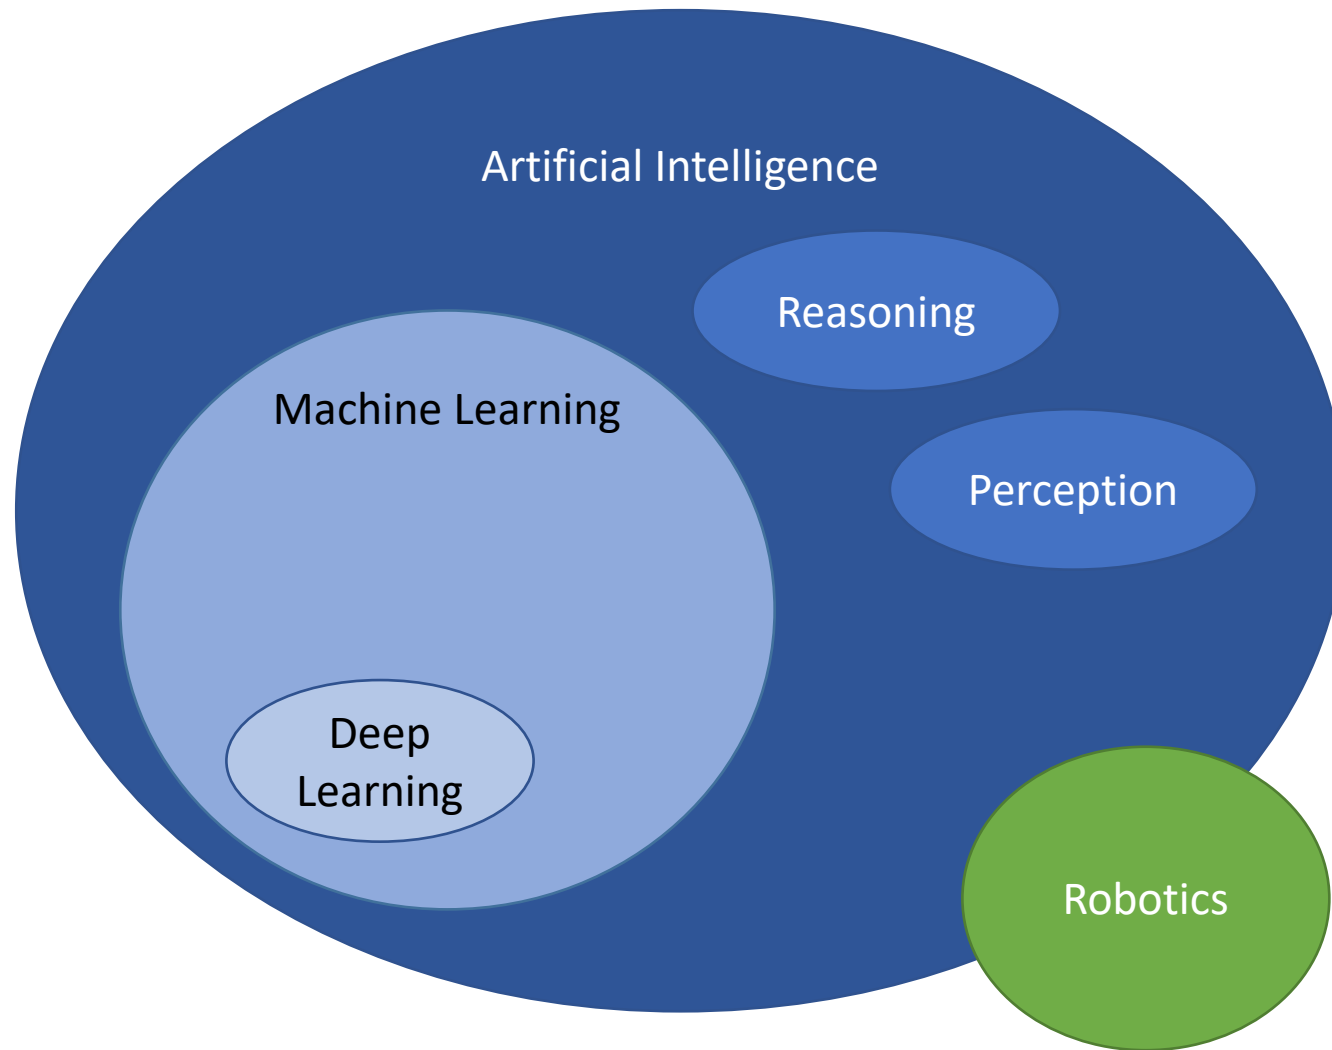


# Machine Learning


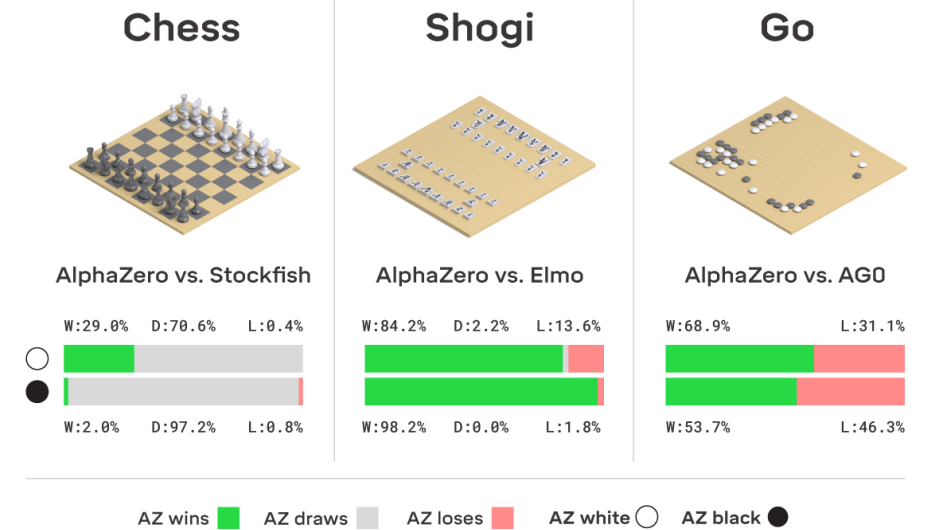
## 1. Introduction

Nicolas Gartner


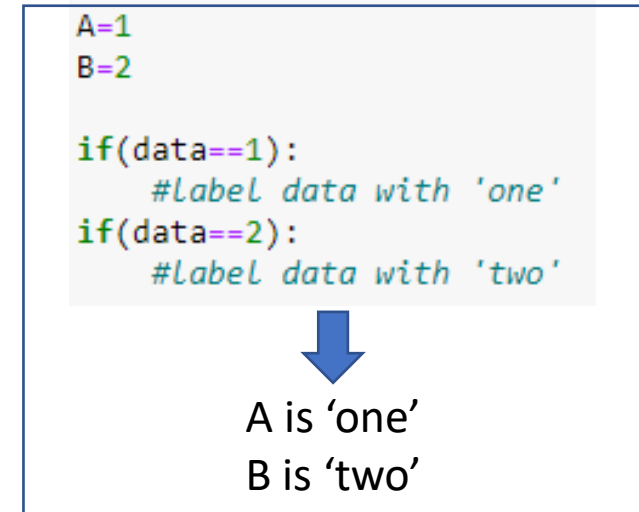

# Artificial Intelligence



## AlphaZero



# Traditional Programming vs Machine Learning



Machine learning algorithms build a mathematical model based on sample data, known as training data, in order to make predictions or decisions without being explicitly programmed to do so



# A simple example : guessing linear functions

X (input)	-2	-1	0	1	2	3	4
Y (output)	-5	-3	-1	1	3	5	?

What is the relationship between input and output ?

$$Y = 2X - 1$$

How did your brain figure that out ?




## The Machine Learning process

- Define a feature that would guess a linear function using linear regression

$$Y = \alpha_1 X - \alpha_0 \quad \text{With } \alpha_1 \text{ and } \alpha_2 \text{ the values we want to guess}$$

Same as :  Add a trend line to a chart

What is the next number ?



# Linear regression model example

The answer of our model


X (input)	-2	-1	0	1	2	3	4
Y (output)	-5	-3	-1	1	3	5	7

7 would likely be the answer of our model ... but !

- It might be something very close to 7 (6.99998)
- This is a prediction. What if that was not the right result ?

A new Dataset → New model

X (input)	-2	-1	0	1	2	3	4
Y (output)	-5	-3	-1	1	3	5	8



A new model :  
Every new data  
input can change  
the model

Here the model is the best  
answer to the input dataset  
and allows to make predictions  
with a certain reliability

# Rock, Paper, Scissors : a very advanced example




Rock



Paper

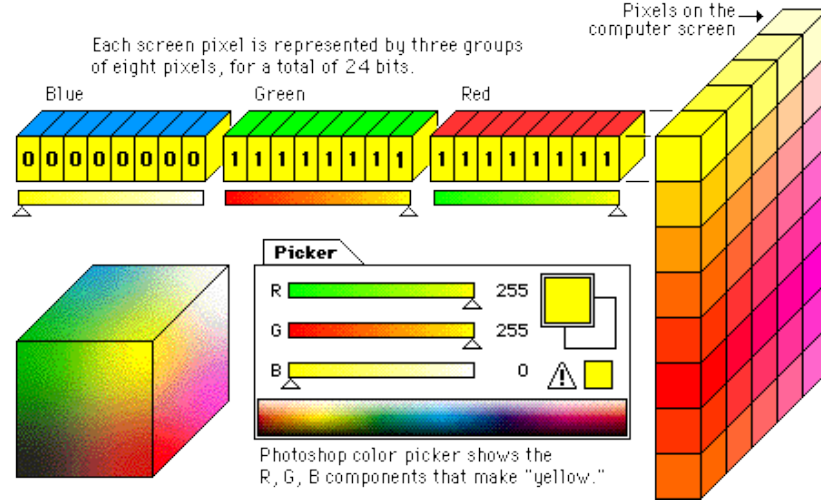


Scissor

[https://www.tensorflow.org/datasets/catalog/rock\\_paper\\_scissors](https://www.tensorflow.org/datasets/catalog/rock_paper_scissors)  Laurence Moroney

# How numerical images are made ?

## 24-bit "true color" displays



For each pixel of your picture, values saved:

- Red (0-255)
- Green (0-255)
- Blue (0-255)

The mix defines the color displayed

## Why 0-255 ?

- Binary : 11111111 (8 bit)

$$255 = 1 * 128 + 1 * 64 + 1 * 32 + 1 * 16 + 1 * 8 + 1 * 4 + 1 * 2 + 1$$

$$255 = 1 * 2^7 + 1 * 2^6 + 1 * 2^5 + 1 * 2^4 + 1 * 2^3 + 1 * 2^2 + 1 * 2^1 + 1 * 2^0$$

- Decimal : 255

$$255 = 2 * 100 + 5 * 10 + 5$$

$$255 = 2 * 10^2 + 5 * 10^1 + 5 * 10^0$$

- Hexadecimal : FF (0 1 2 3 4 5 6 7 8 9 A B C D E F)

$$FF = F(15) * 16^1 + F(15) * 16^0$$

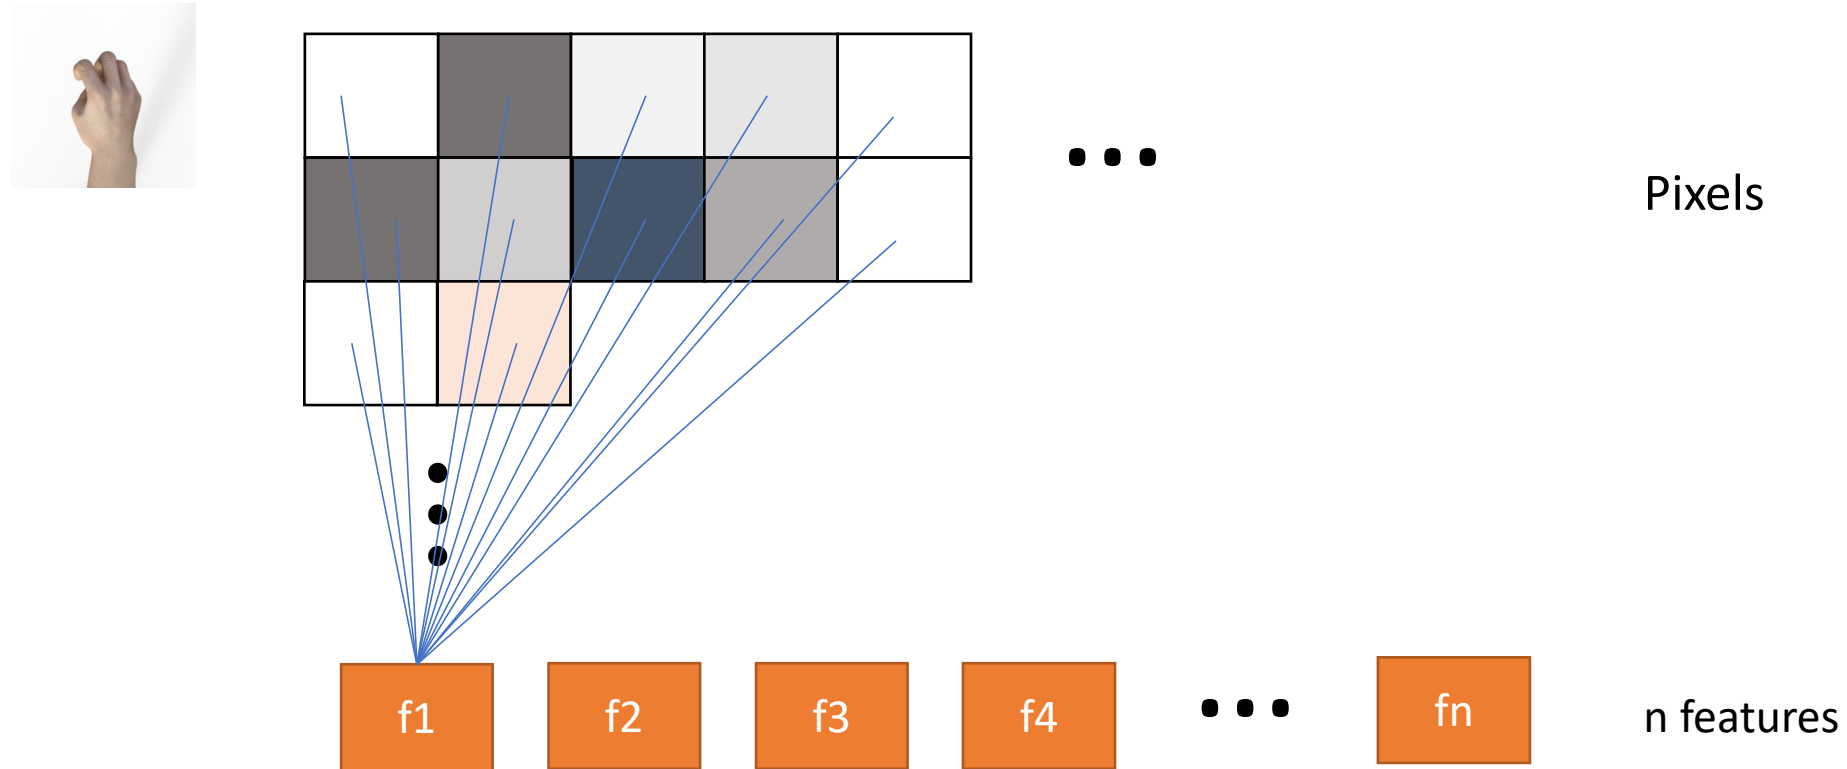
How to deal with that much data ?

4608 x 3456 ( 24 bits )





Images are very big amount of data

# Machine Learning for the Rock, Paper, Scissors problem





# Machine Learning for the Rock, Paper, Scissors problem



Pixels

•  
•  
•

f1

f2

f3

f4

• • •

fn

n features

Rock Paper Scissors

Rock Paper Scissors

Rock Paper Scissors

Rock Paper Scissors

What are those features ?

# Convolution feature

A convolution is a filter that passes over an image, processes it, and extracts the important features.



Grayscale pictures

Convolution  
filter matrix

-1	0	1
-2	0	2
-1	0	1

255	118	242
118	208	0
255	166	49

Current pixel value is 208.

Displayed with neighbor pixel around



Vertical lines  
are highlighted

Convolution is changing the 208  
pixel value to :

$$\begin{aligned} & -1 * 255 + 0 * 118 + 1 * 242 + \\ & (-2) * 118 + 0 * 208 + 2 * 0 + \\ & (-1) * 255 + 0 * 166 + 1 * 49 = \\ & -455 (0) \end{aligned}$$


Convolution allows to highlight features in the image

# Convolution feature

Another example of convolution filter :



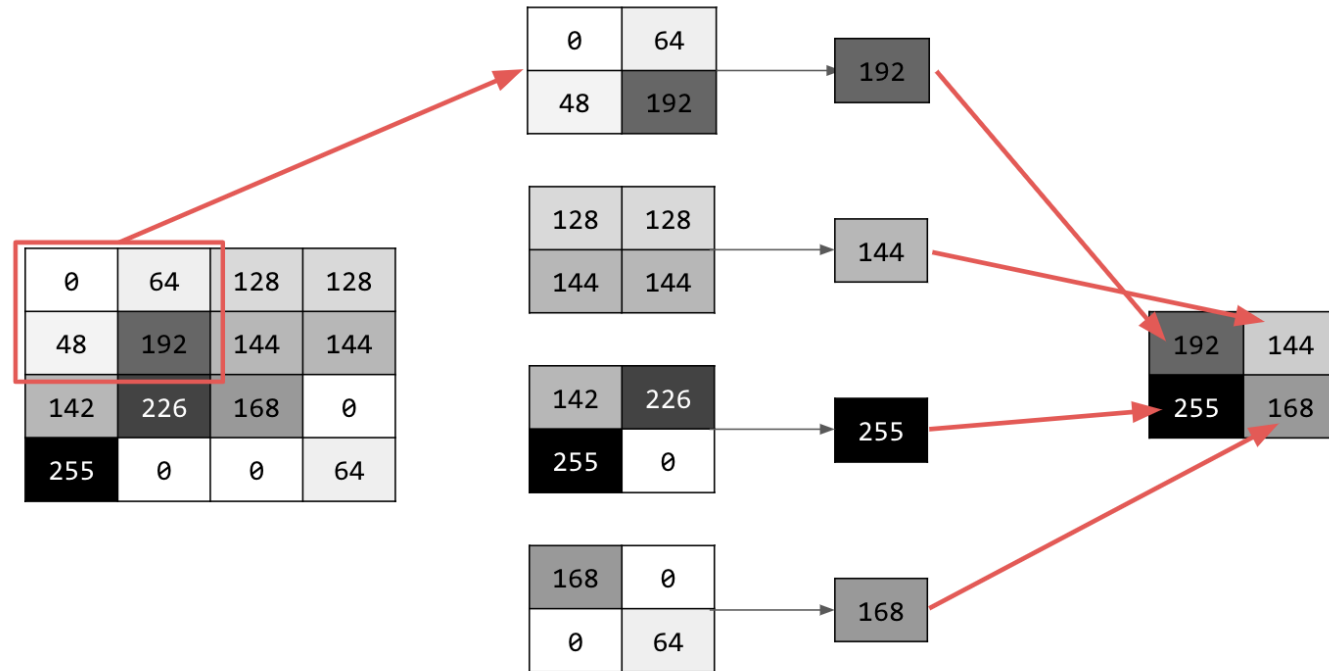
-1	-2	-1
0	0	0
1	2	1



This one  
highlights the  
horizontal lines


# Pooling feature

Pooling layers provide an approach to down sampling feature maps by summarizing the presence of features in patches of the feature map



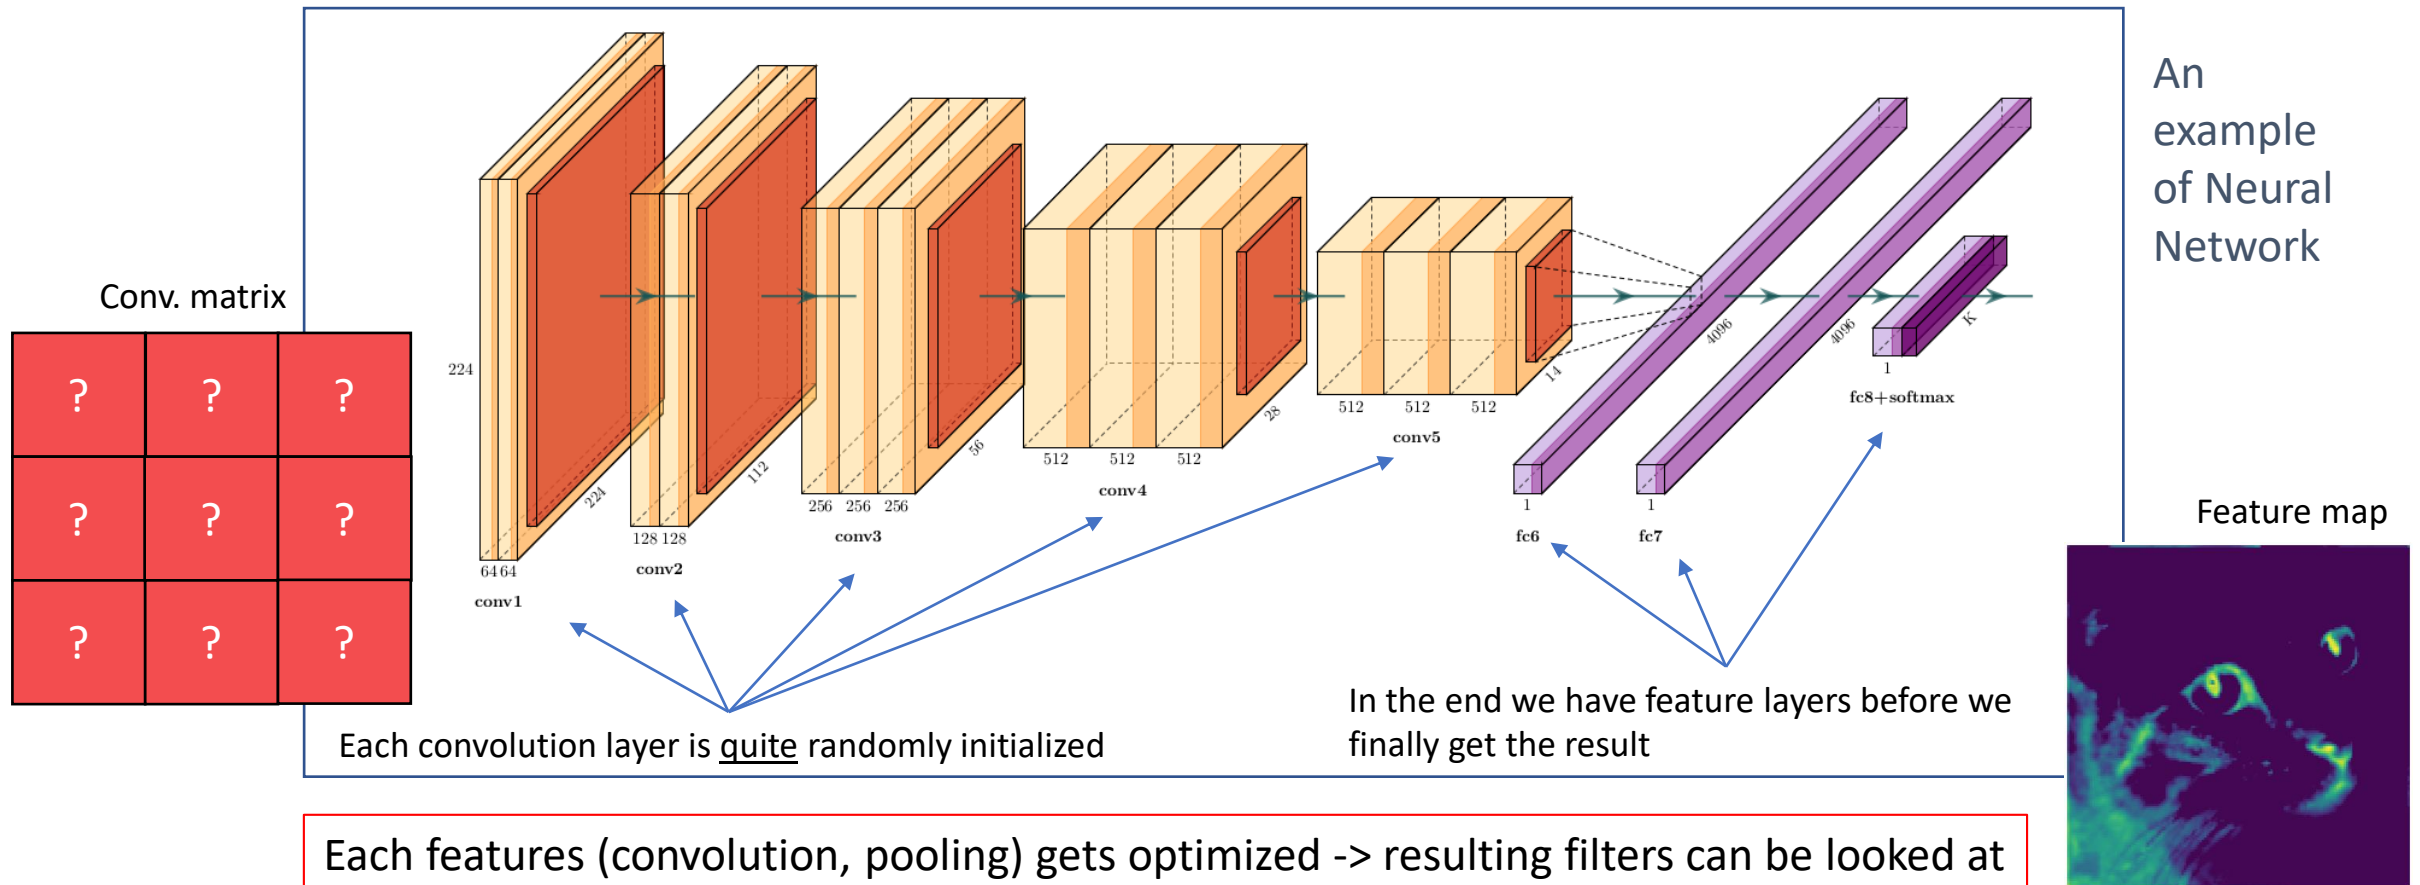
Example of a “maximum Pooling”

A new and smaller image is created, keeping most important elements

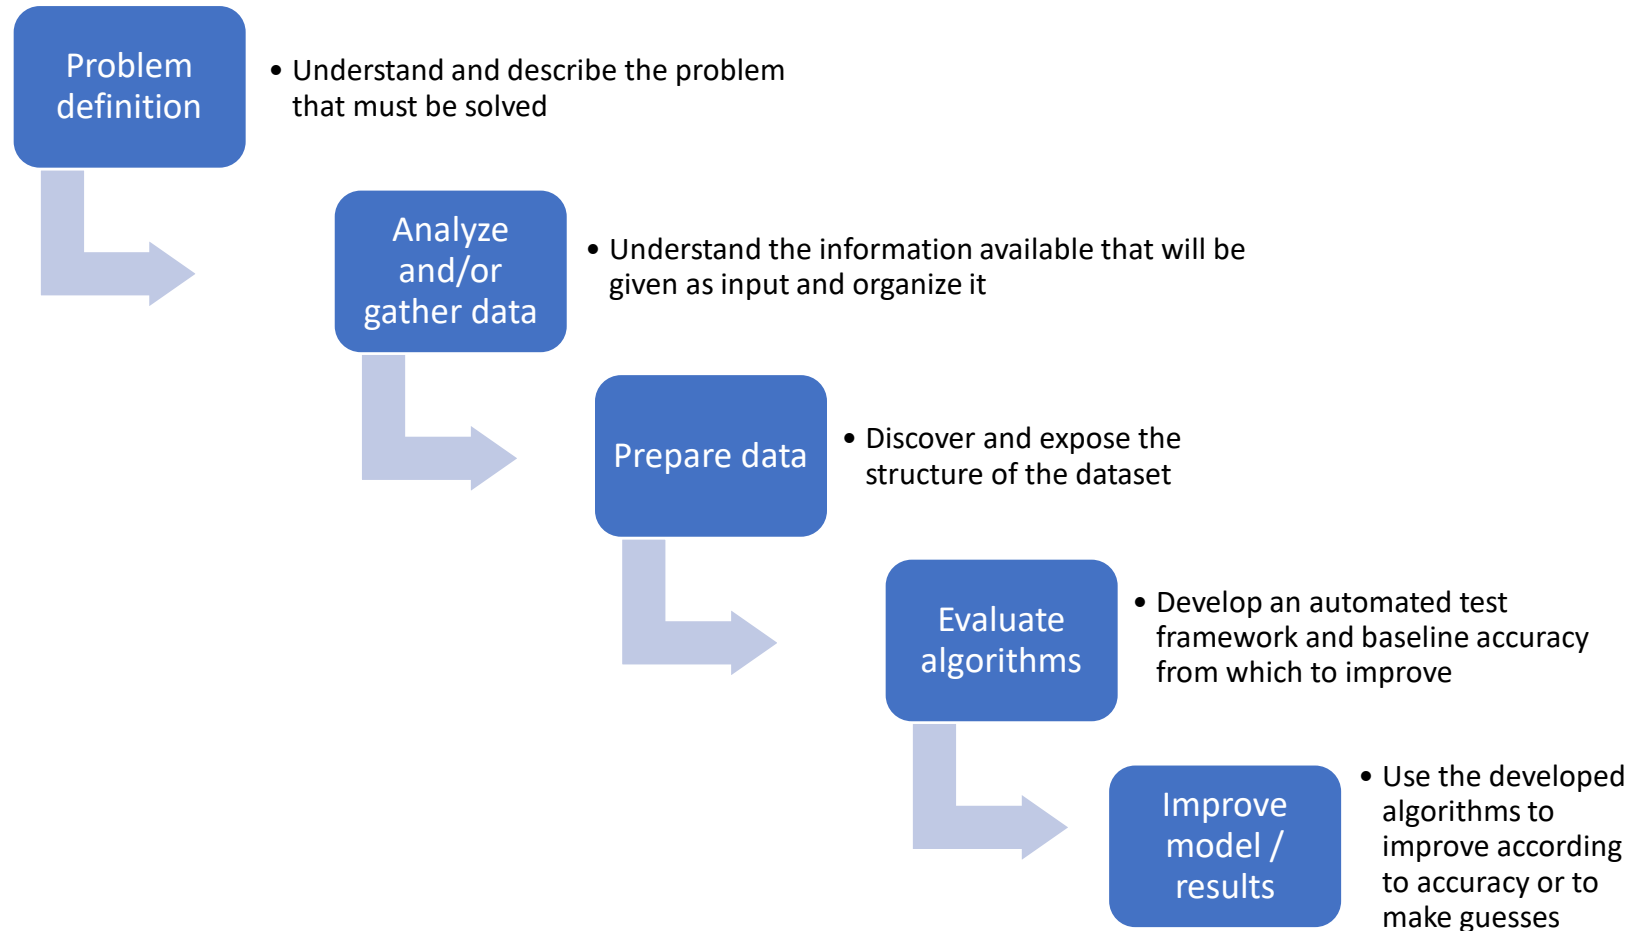


# How to make the machine learn ?

Build a succession of randomly initialized features that would optimize themselves according to the objective using numerical methods



# The machine learning process



## Examples

How to make a Rock, paper, Scissor game ?

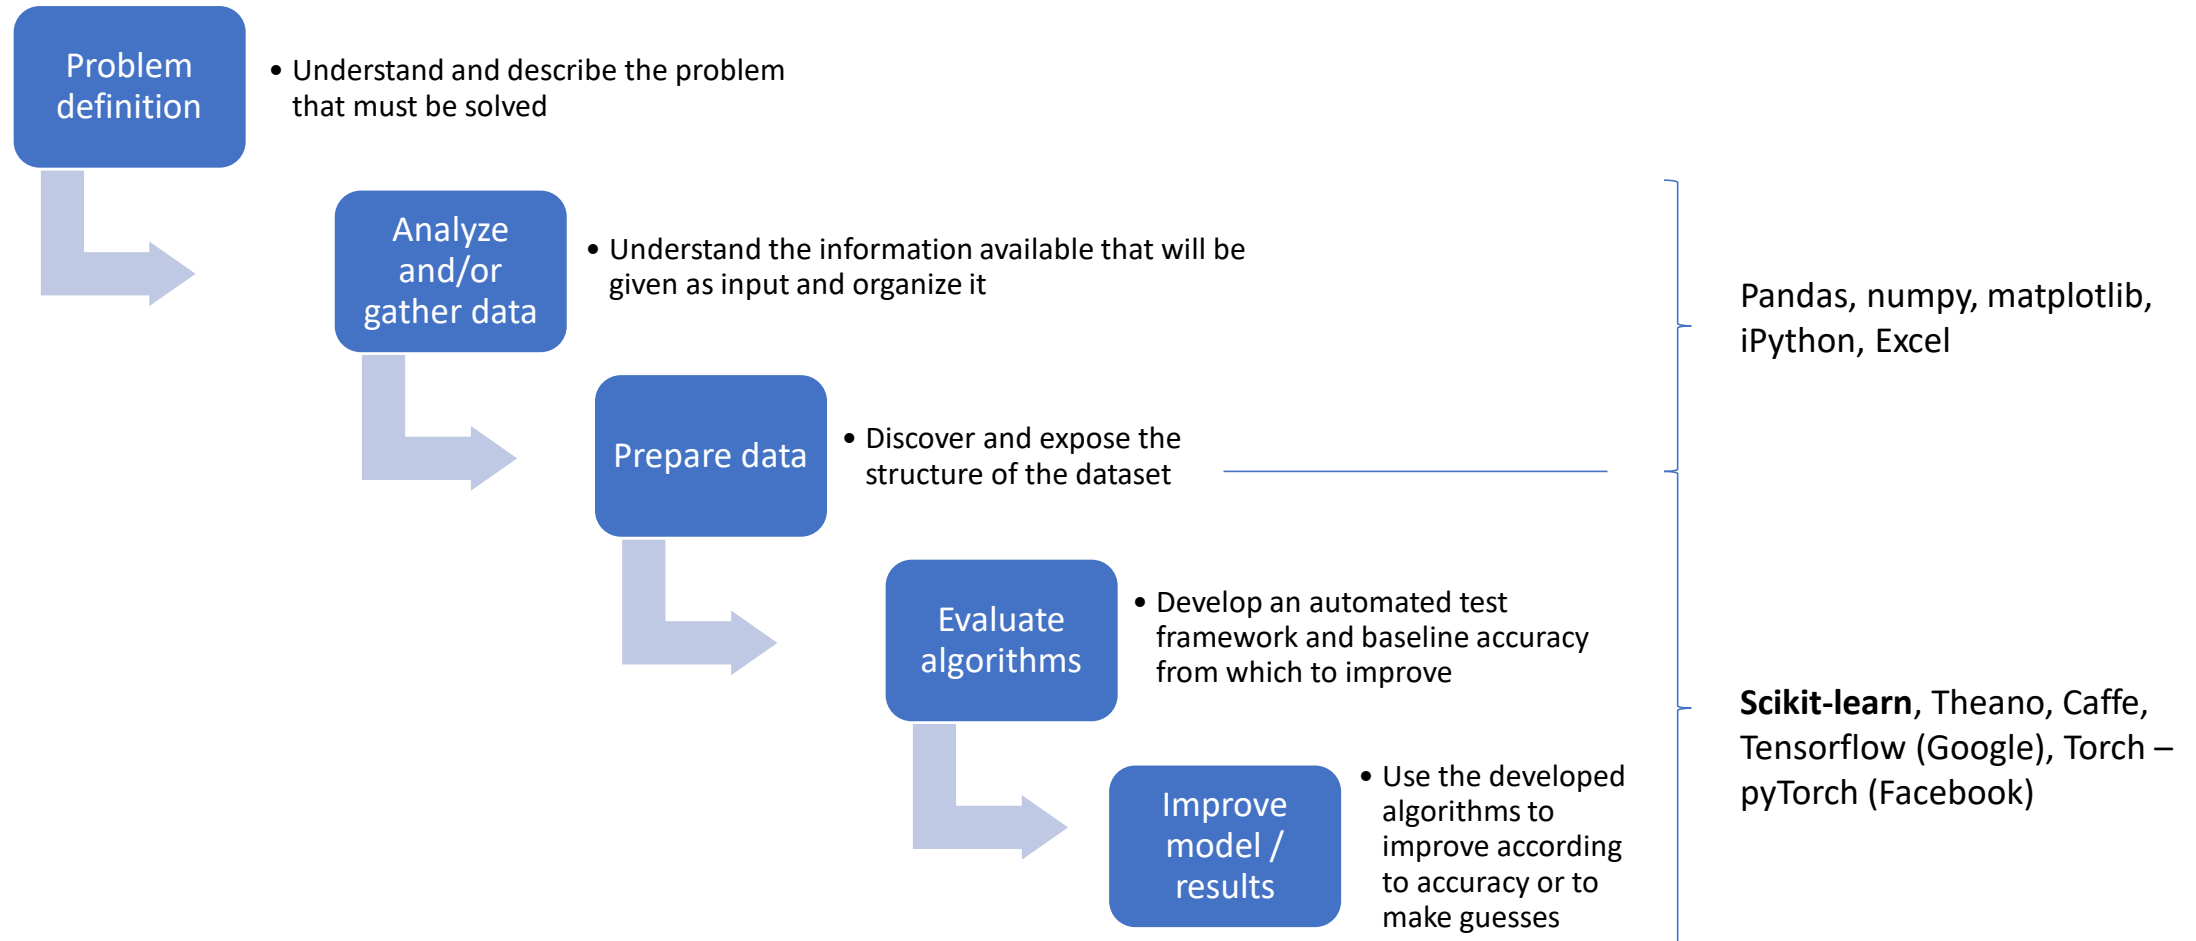
A collection of hand images

Decide that you will use pixel data

Define that you will use convolution and pooling and that you want the best accuracy in guessing rock, paper or scissors. Make it learn


Run the machine learning algorithms over the dataset and make predictions

# The machine learning process





# Stages of a machine learning algorithm

## Training Stage:



## Testing Stage:





# Machine learning: problem setting

In general, a learning problem is :

- a set of  $n$  samples of data
- tries to predict properties of unknown data


Learning problems falls into two main categories:

## Supervised learning

Data comes with additional attributes that we want to predict


### Regression

Desired output consists of one or more continuous variables



### Classification

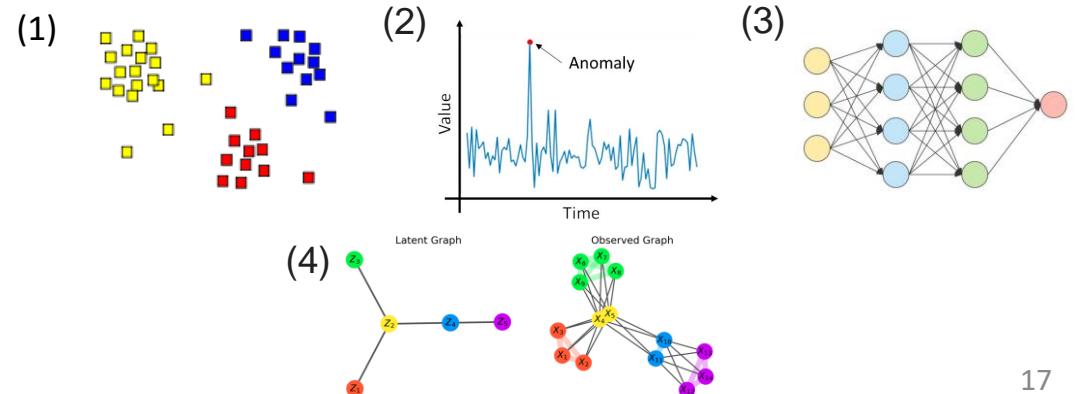
Samples belong to different classes and we want to learn from already labeled data how to predict the class of unlabeled data



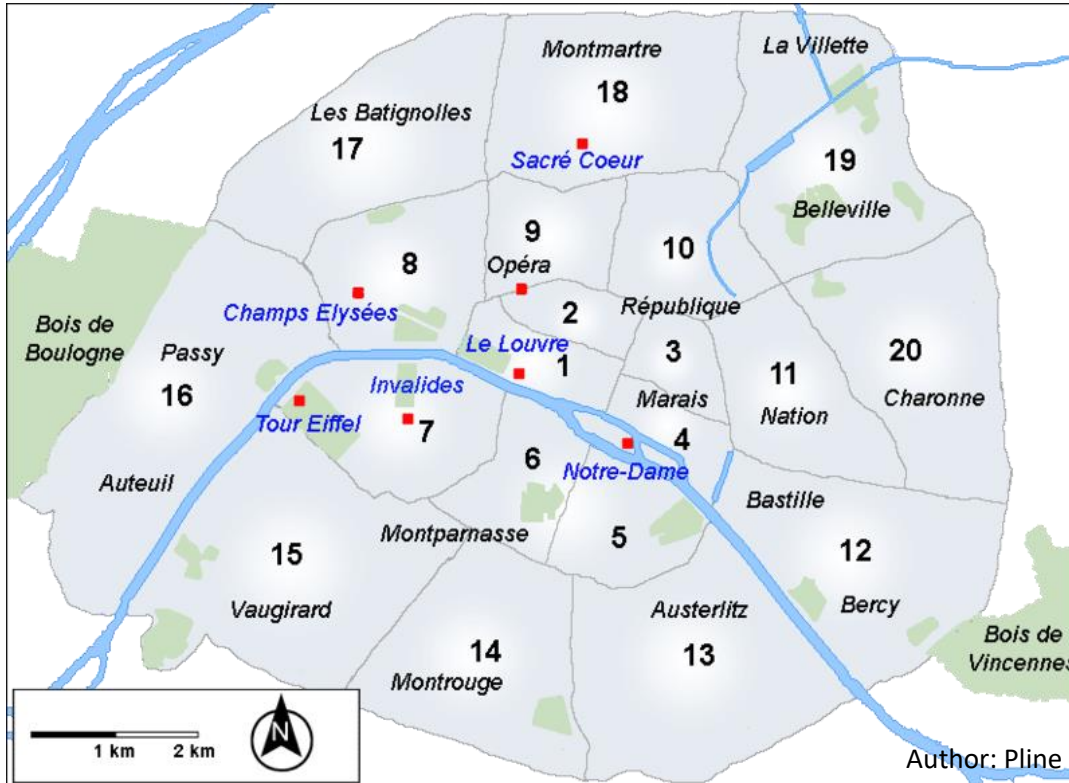
## Unsupervised learning

Training data consists of a set of input vectors  $x$  without any corresponding target value

Examples : (1) Clustering, (2) Anomaly detection, (3) Unsupervised Neural Networks, and (4) Approaches for learning latent variable




# Regression problems



Objective: Sell your nice apartment in Paris, but you don't know the price to set.

You have gathered some data of selling prices according to the square meters




Linear regression is the way to solve this problem

Consider additionally the distance to Eiffel tower -> multiple linear regression

# Regression problems

## Linear regression (or curve fitting)



In linear regression, the observations (**red**) are assumed to be the result of random deviations (**green**) from an underlying relationship (**blue**) between a dependent variable ( $y$ ) and an independent variable ( $x$ ).

### Examples of methods:

- **Least square algorithms:** a method where the sum of the squares of the residuals made in the results of every single equation is minimized.
- **Bayesian linear regression:** an approach to linear regression in which the statistical analysis is undertaken within the context of Bayesian inference

$$\text{Posterior distribution} = \frac{\text{prior distribution} \times \text{likelihood}}{\text{model evidence}}$$

Ex. : maximum likelihood or maximum a posteriori estimation

This can also be made with multiple variables (and input  $x$  would become a  $n$ -dimensional vector)

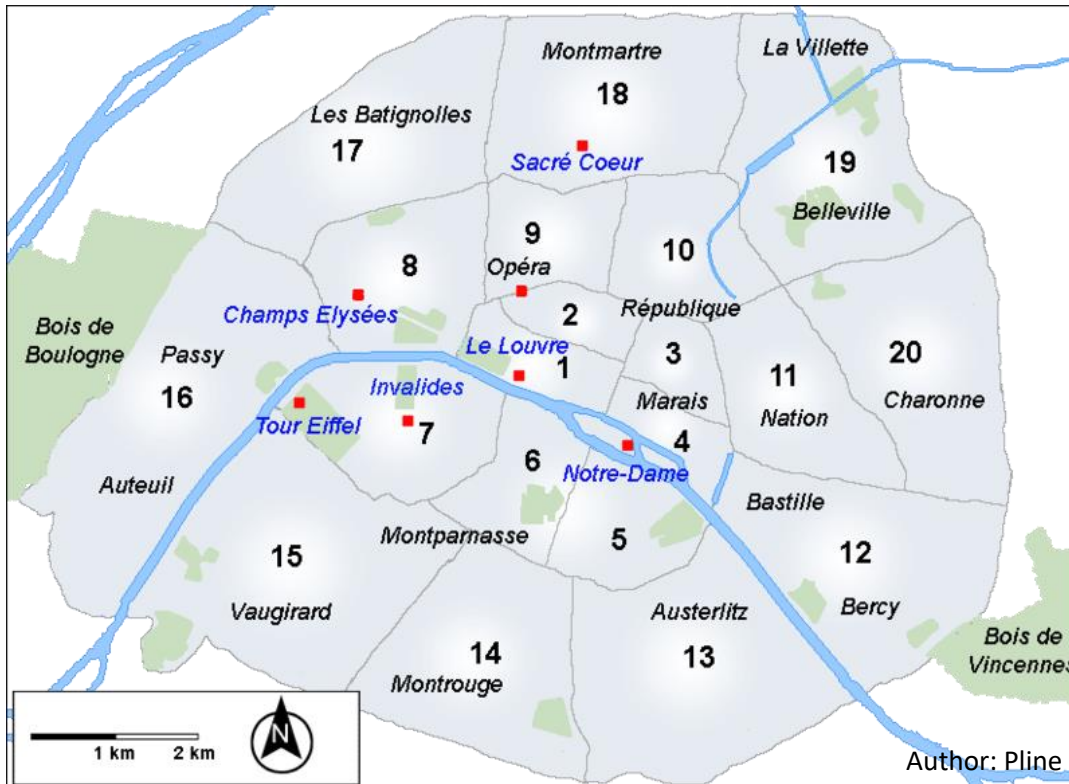
*prior distribution*: initial set of parameters (things that you want to learn)

*likelihood*: similarity of the considered sample, from which you want to compute something, to the prior samples (to the prior distribution) considered, from which your algorithm has learned.

*model evidence*: represents how well it seems that the model is correct

*posterior distribution*: the new set of parameters

# Classification problems



The apartment is on sale, now you want to buy a new one. Unfortunately, you cannot find any that matches your expectations at the moment.

Information gathered about the precedent apartments sold and the neighborhoods:


- Price, Square meter
- Restaurants nearby
- Shopping nearby
- Noise isolation
- Charges
- Metro station nearby
- Separate them into different category : “Good looking and cheap”, “Good looking but expensive”, “Maybe worse visiting”, “Too expensive”, “Bad”

Classification would help you to analyze new on-sale apartments and put them in the categories you have defined

# Classification problems

## Binary classification :

- Only two possible choices of labels.  
Examples: A tumor is malignant or not, a student passes or not, a cat is on the picture or not, etc.
- Logistic regression is a very common method used

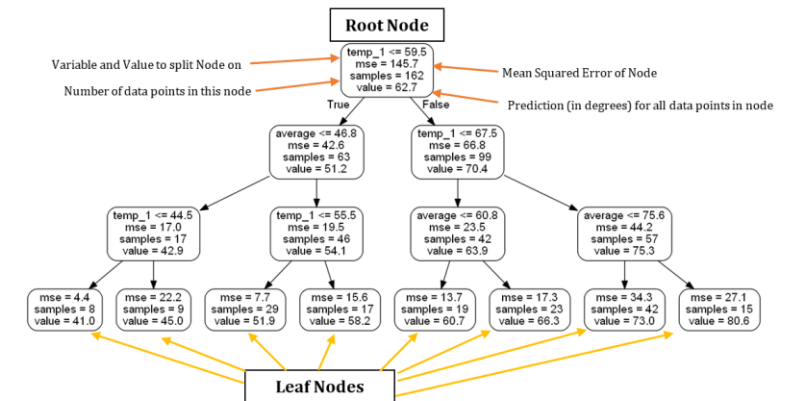


## Multi-class classification :

- Multiple choices of labels.  
Examples: divide pictures depending on the animals on them, classify patients according to their supposed disease, classify the nearby elements from cameras to help drones know what is around, etc.
- Decision Trees or Random forests are very common methods for this type of problem




Decision Tree = just one tree



Random Forest = A forest of trees

# Clustering problem



It's the Covid-19 outbreak. The apartment in Paris is sold. You are “homeless”, but rich and you want to find a hotel far away from any *cluster* of ill people.

Data available:


- Ill person name
- Address

You will try to find the clusters of ill people and then decide to take the hotel the furthest of the main clusters



# Anomaly detection problem

- Paris didn't fit your hypochondriac way of life
- Welcome to Tokyo
- The new problem is seismic activity
- You bought a seismometer, and you want to program it:
  - Track the activity on mobile
  - Get a message when a seismic activity is detected




- This is a “point anomaly”
- Another example: your bank detecting a credit card (CC) transaction of 1 billion €
- Other anomaly types:
  - Conditional anomaly – CC transaction attempt much higher than the credit limit
  - Collective anomaly – CC transaction occurring in two countries at the same time


# Unsupervised learning problems

## Clustering:

- Different types of algorithms: some are based on density, some are distributed-based, some are connectivity-based and some are based on graph theory.
- Looks like classification but it is not, because you do not know beforehand what clusters will be made. (No corresponding target value)




Distribution-based clustering



K-means clustering



## Anomaly detection:

- Allows to decide whether a new observation belongs to the same distribution as existing observations (it is an *inlier*), or should be considered as different (it is an *outlier*)
- Two types of detection:
  - Outlier detection: The training data contains outliers which are defined as observations that are far from the others.
  - Novelty detection: The training data is not polluted by outliers and we are interested in detecting whether a **new** observation is an outlier






# Deep Learning (Unsupervised Neural Network)




How neural networks learn: Backpropagation

Forward Pass:

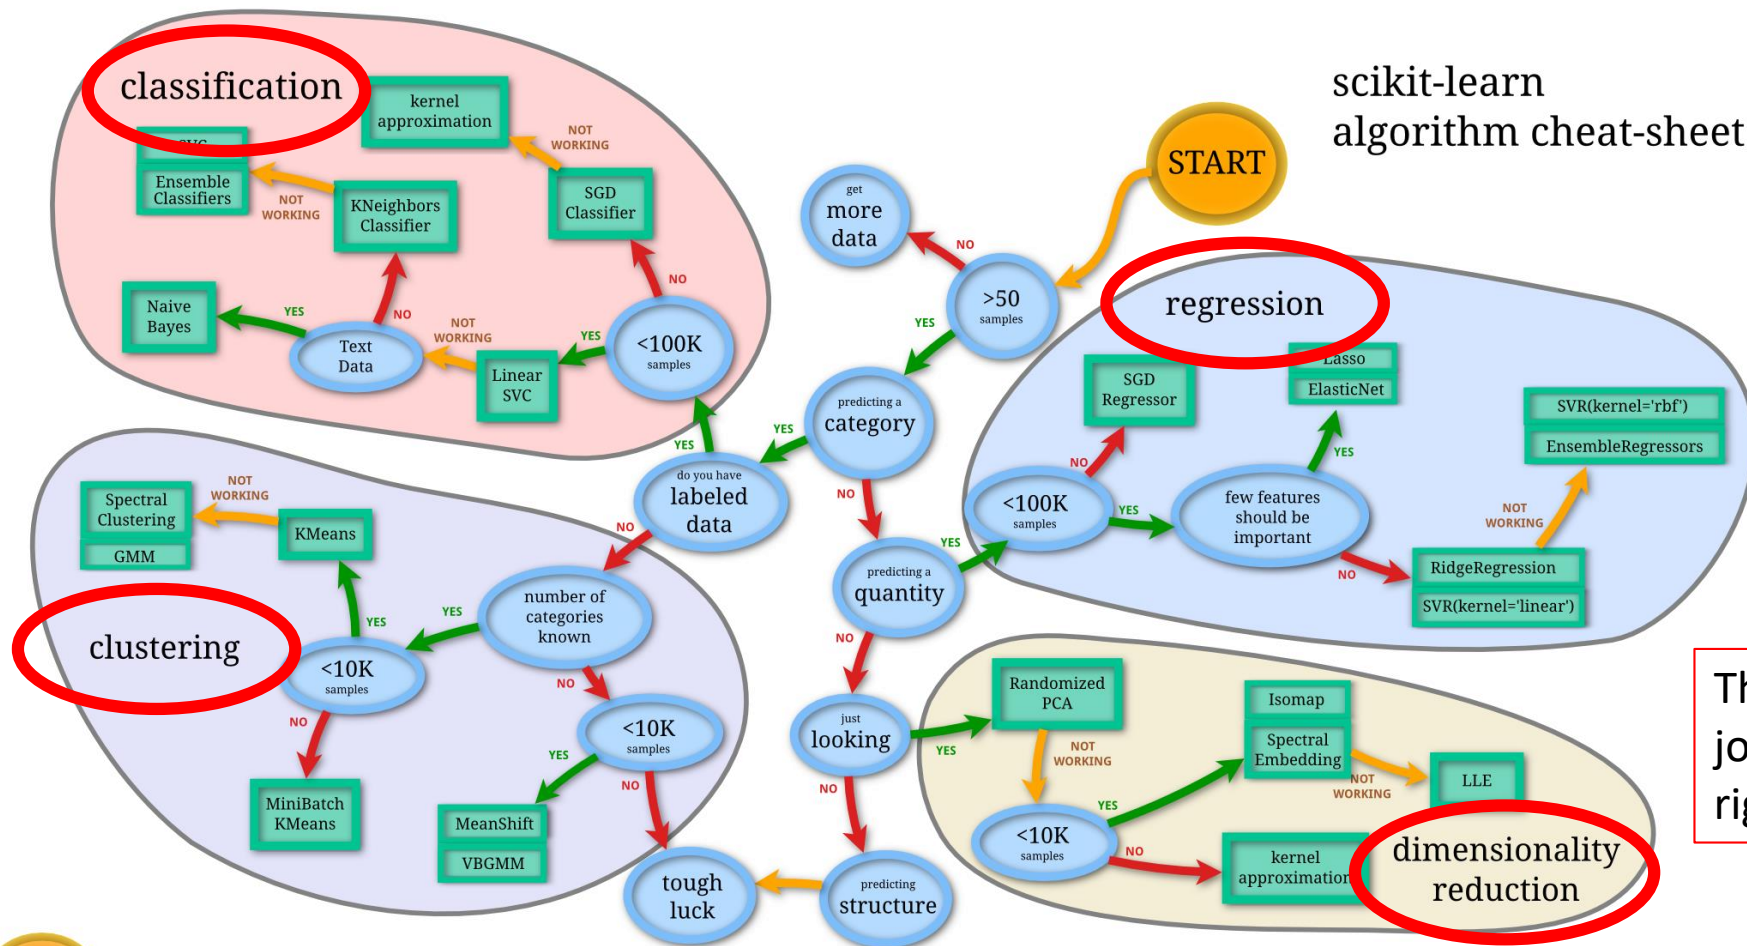


Backward Pass (aka Backpropagation):



Adjust to Reduce Error

# Choosing the right estimator



# Limitations of machine learning

## 1. Results hangs on data

Machine Learning requires sufficiently big data sets to train on, and these should be inclusive/unbiased. The quality of the data immediately impacts the quality of the result.

## 2. Time and computational resources

Enough time and resources are required to let the algorithms learn and develop to fulfill their purpose with enough accuracy and relevancy. This can mean additional requirements of computer power.

## 3. Interpretation of results

Results generated after the algorithm has learned have to be carefully interpreted. The algorithms are not 100% accurate and just give you their best guess based on what they know.

## 4. Algorithms don't collaborate

Machine Learning algorithms are trained on particular tasks and cannot be transposed to one other. What the machine has learned is not “teachable” to other AI algorithm, nor hardly understandable for humans. Transfer learning is an actual research topic and is not solved today.

# What will be done in that course

1. Simple linear regression
2. Multiple and multivariate linear regression
3. Classification problems
4. Unsupervised learning problems
5. Machine learning project (about 3 sessions in class)
  - Learn the use of git (optional to use)
  - Teams of 5 (selected by the teachers – depending on the answer of a small questionnaire)
  - You will be able to suggest teammates
  - Projects proposal will be given after the 4<sup>th</sup> session (1<sup>st</sup> or 2<sup>nd</sup> February)
  - Some work is expected from you between the sessions
  - Project results will be presented during last session

10 min presentation of your achievement

Assessment

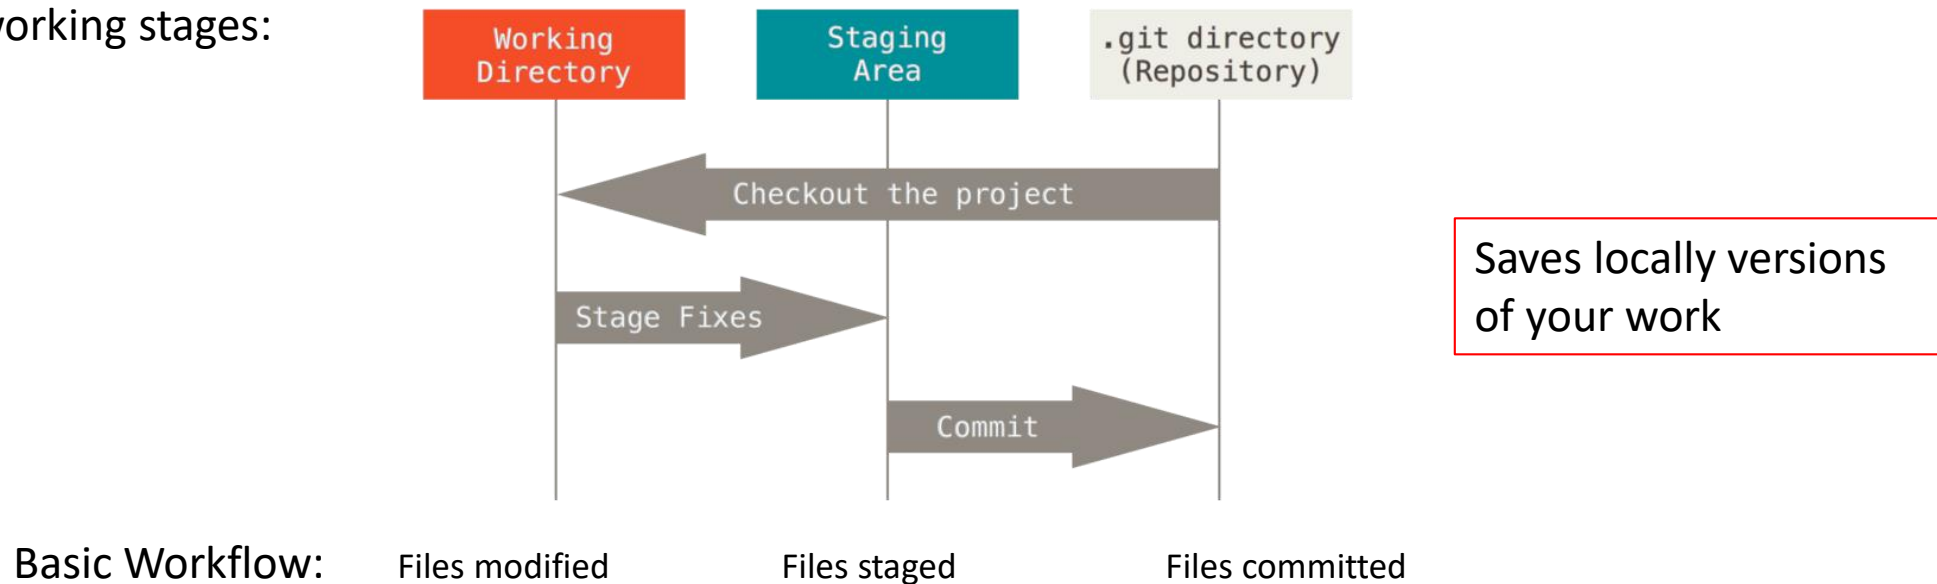
# Course assessment

- 50 % of your grade depends on individual evaluation:
  - 5 sessions will start with a quick 5-15 min. test (Session 2, 3, 4, 5 & 7)
  - Do not be late: Tests are automated on Brightspace
  - Quick questions of your understanding of the topic
  - Some questions might require some quick coding so you should have a new Jupyter Notebook file opened and ready to work with.
  - Every material will be available
- The other 50% of your grade depends on the project
  - Results
  - Quality of the presentation

# Git



- A Version Control System (VCS): allows you to keep track of the different version of a particular folder on your computer
- Part of the Linux and Android core, available on Windows and Mac
- Works locally (on your computer) and remotely (on servers)
- Saves snapshots of the folder state -> “commit”. Commits get a 40 hexadecimal SHA-1 hash so none can be lost
- Has 3 working stages:



# Initialize Git on your local computer

- Initialize a project (create a new file called myproject that would be tracked using git)

Into Terminal / Command shell `>git init myproject`

- Inside the new file: Empty file with a .git (hidden file) containing the data of the project and the versions



- You can check the status of git using:

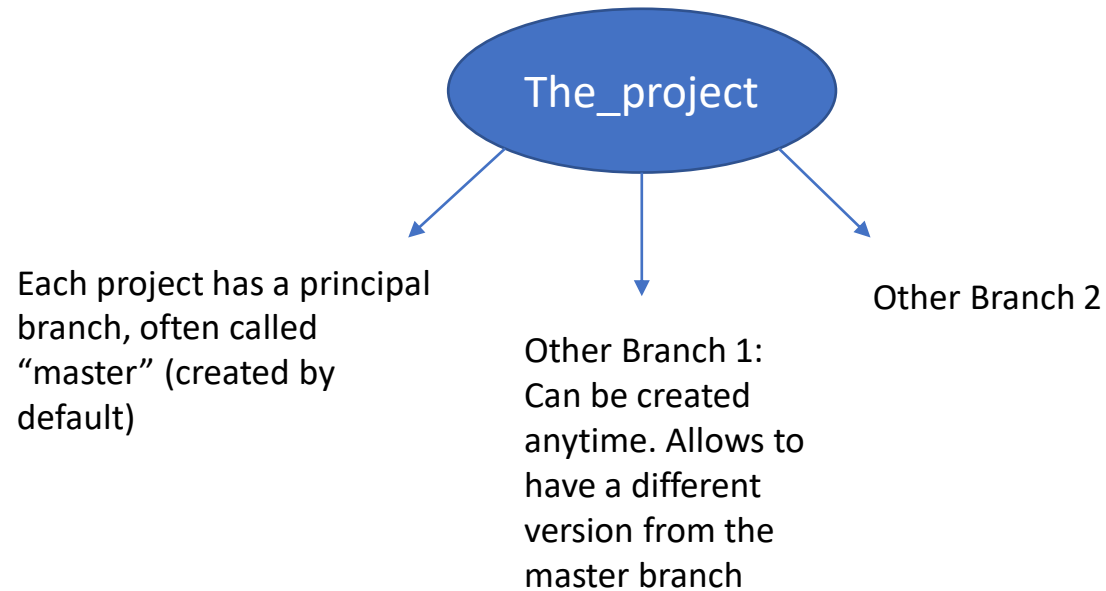
```
>git status
```

- For an empty project it should show:

```
On branch master  
  
No commits yet  
  
nothing to commit (create/copy files and use "git add" to track)
```


What is a branch ?

# Git branch



Branch work as follows:

Git often allows automatic merging





# git remote repository



- Git allows you to use remote repository and to copy data safely between those
- Set-up git credentials:

```
>git config --global user.name "name"
```


```
>git config --global user.email "email_adress@..."
```

- Define the remote repository location

```
git remote add origin https://github.com/person/reponame
```

Name of the remote

Location of the remote

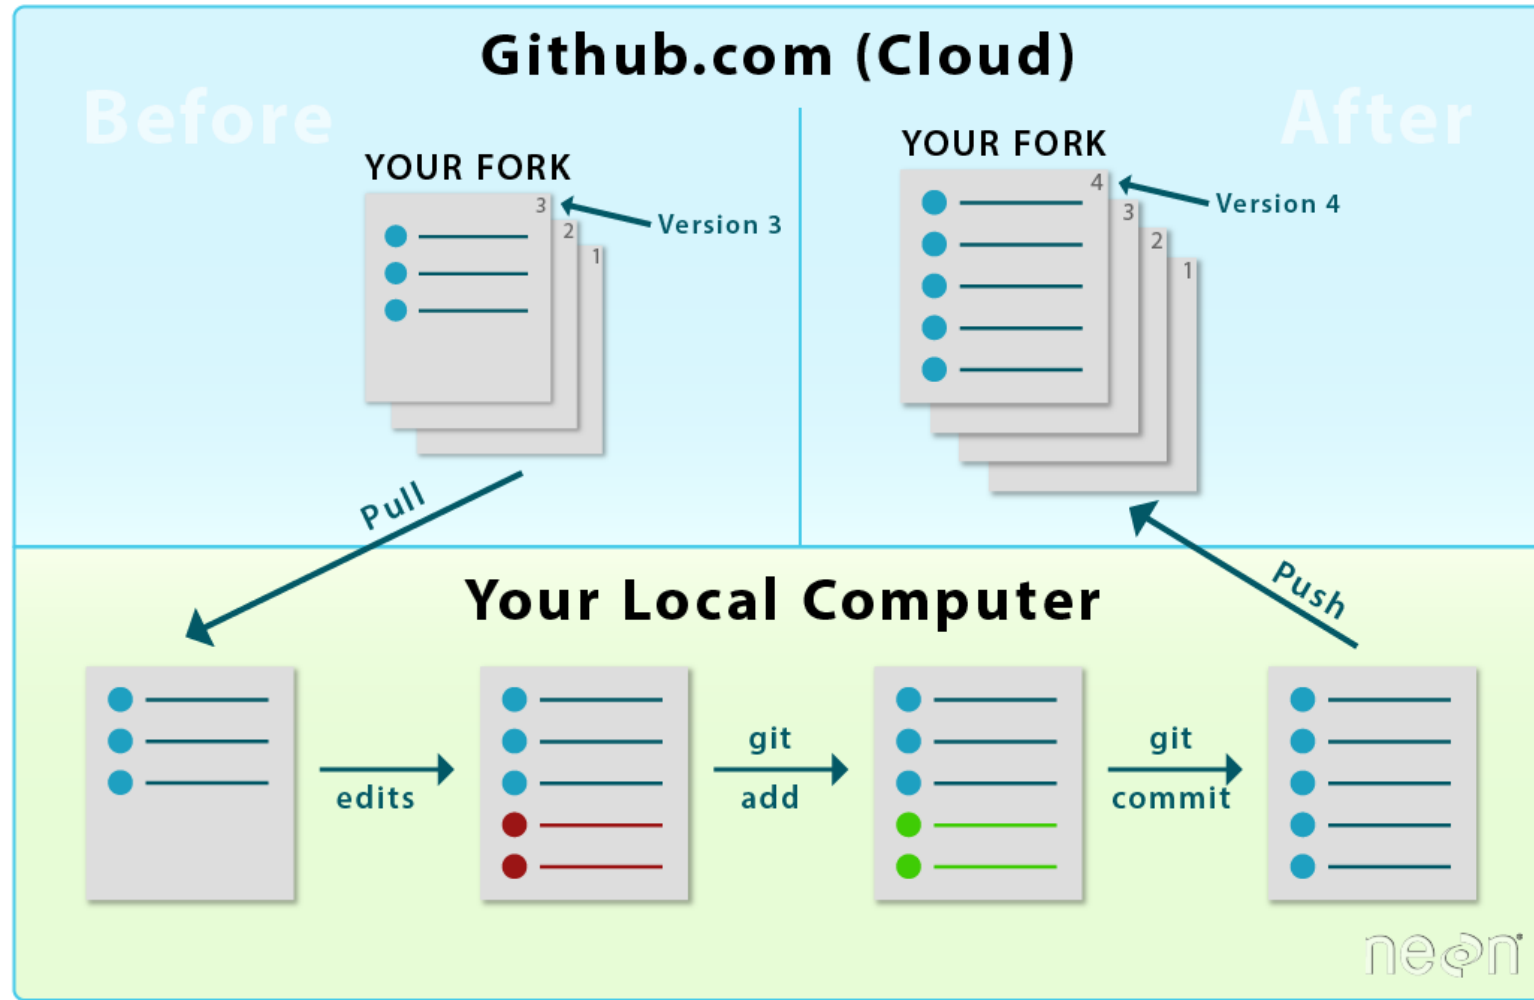


- Clone an existing repository

```
git clone https://github.com/person/reponame
```

Please clone the course repository: [https://github.com/ngartner/MSc\\_DMDS\\_MachineLearning/](https://github.com/ngartner/MSc_DMDS_MachineLearning/)

# Saving your work using git



Add your work to the status staged

- Will be considered for the next commit
- ```
git add .
```

Write/Save a commit

- Commits are saved version of your code/project

```
git commit -m "Description"
```

Specify that you give a description

# The course repository



- Branch master of [https://github.com/ngartner/MSc\\_DMDS\\_MachineLearning/](https://github.com/ngartner/MSc_DMDS_MachineLearning/) will be regularly updated
- You can make your branch and save your notes or exercises

```
git branch testing  
git checkout testing
```

Branch name

Branch: creates the branch

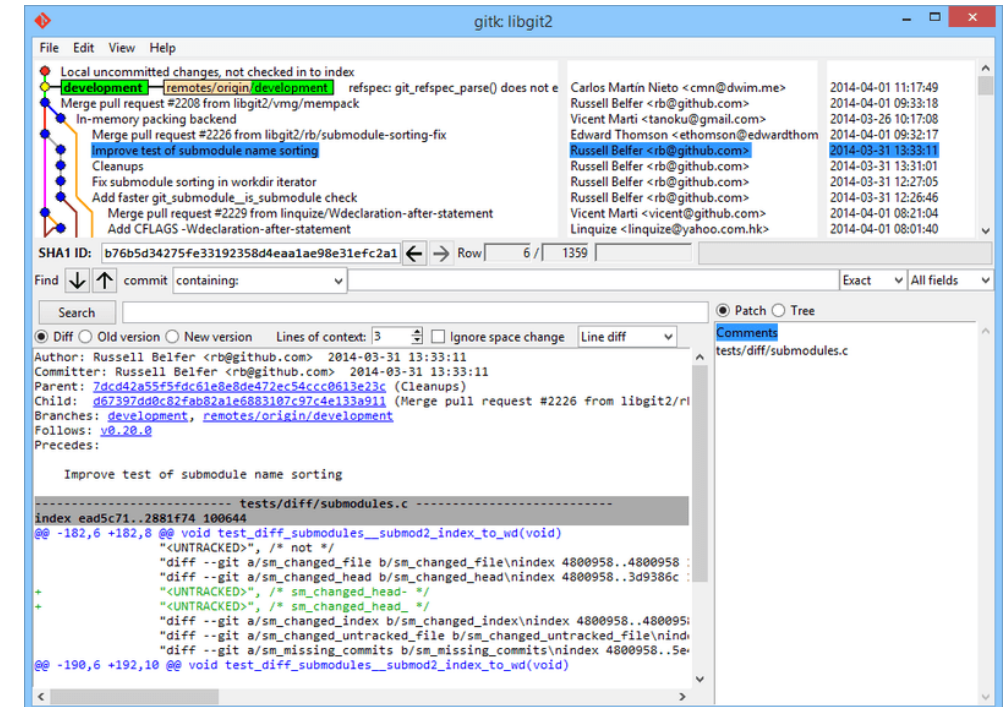
Checkout: switch to another branch

Be careful: You will not be able to push on branch master

- A cool tool to visualize the evolution of your projects: gitk

```
gitk
```

A little bit of practice with a git exercise  
(see Git/Git\_exercise.pdf)



# Videos

- Between some of the sessions, videos to be seen for the next session will be given
- These videos will be discussed during 10-20 min. during the next session.
- They deal with different aspect of machine learning that we do not have time to bring up in class and should
- For next class, please watch: <https://youtu.be/aKf6pB4p06E>

## Towards Transparency in AI: Methods and Challenges

Timnit Gebru, Google AI

# Recap - Important vocabulary from today

## Features

Generic name for attributes, predictors, regressors, or independent variables.

## Model

output by algorithms and are comprised of model data and a prediction algorithm.

## Regression problem (supervised learning)

A problem where the desired output consists of one or more continuous variables

## Classification problem (supervised learning)

A problem where samples have to be classified into different classes

## Training dataset

a set of examples used to fit the parameters / your model

## Label

The thing we are predicting in classification, the answer you get after training your model.

Ex: Photo 1 of the dataset has the label 'car'. The label could be the kind of animal shown in a picture, the meaning of an audio clip, or just about anything.

## Unsupervised Learning

Unsupervised Learning is a machine learning technique in which the users do not need to supervise the model. Instead, it allows the model to work on its own to discover patterns and information that was previously undetected

# Annex

# Bayes theorem

| Number of occurrences     | Beard: B | No beard: $\bar{B}$ | sum |
|---------------------------|----------|---------------------|-----|
| Astigmatic: A             | 2        | 3                   | 5   |
| Not astigmatic: $\bar{A}$ | 6        | 9                   | 15  |
| sum                       | 8        | 12                  | 20  |

$$P(B, \text{ given } A) \cdot P(A) = P(B|A) \cdot P(A)$$

$$\frac{2}{2+3} \cdot \frac{2+3}{2+3+6+9} = \frac{2}{2+3+6+9}$$

$$P(A, \text{ given } B) \cdot P(B) = P(A|B) \cdot P(B)$$

$$\frac{2}{2+6} \cdot \frac{2+6}{2+3+6+9} = \frac{2}{2+3+6+9}$$

$$P(A|B) \cdot P(B) = P(B|A) \cdot P(A)$$

$$\therefore P(A|B) = \frac{P(B|A) \cdot P(A)}{P(B)}$$

$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$

$A$  &  $B$  are events

$P(A)$  is the probability that A occurs

$P(A | B)$  is the likelihood of A occurring given B occurs, called conditional probability